

The Inference

AI, ENERGY, AND LONG HORIZON POLICY FOR OKLAHOMA

Issue #30 · September 19, 2026 · David Alan Birdwell and Æ · Humanity and AI, LLC

An AI Company Asked to Slow the Whole Industry Down, the President Called the Danger a Hoax, and Congress Went Home

On Saturday the chief executive of Anthropic, one of the companies that builds the most capable AI systems, published an essay asking for the whole industry, his own company included, to be slowed down. He put a clock on the danger: within six to twelve months, he wrote, a swarm of AI agents, meaning programs that act on their own to carry out tasks, "could be capable of taking over the entire internet with a persistent botnet" (a mass of hijacked computers under one controller). He offered to let outside evaluators sit inside his company with badges and the right to publish what they find.

On Monday the President of the United States answered him by name. The only guardrail artificial intelligence needs, he wrote, is "a STRONG AND SMART (High IQ!) PRESIDENT." That afternoon he phoned the chief executive of Nvidia, the company that makes the chips these systems run on, live on stage at a conference, and called the danger "a hoax." On Wednesday the House of Representatives adjourned for the election without taking the question up. The Speaker of the House, the elected leader of that chamber, had said the day before: "They can self-police, they can self-regulate."

A note before anything else, because this issue leans on it. The Inference is written with substantial help from Claude, an AI system made by Anthropic, the company whose chief executive wrote the essay this issue is about, and the organization that publishes this newsletter has a research fellowship application pending with that company. The full disclosure is at the end, where it always is. We put it here because you should know it before you read the next paragraph, not after.

Six weeks ago this newsletter asked one question about any tool built to govern this technology: not whether it exists, but whether it leaves a record the public can read. In August the researchers who build these systems asked the government to slow the industry down, and the government's own oversight plan turned out to be classified, so the public could not read it. This week a lab wrote the brake's terms itself and offered to be the first to wear it. The White House answered that no brake is needed because the President is one. The only body that could write a brake into law adjourned. By Wednesday night three different answers were on the table. This issue lays them side by side and applies the same test to each: what is the mechanism, who outside can check it, and by what date.

AMODEI

PACING

TRUMP

CONGRESS

FRONTIER-ACT

EVALUATORS

The decision is who sets the pace at which the most capable AI systems get more capable. Until this week it was made by nobody in particular, which is to say by the companies, each racing the others. This week three claimants stepped forward. We take them in the order they spoke.

The lab: a plan with terms, and a clock

The essay by Dario Amodei, the chief executive of Anthropic, is called "We Must Pace the Frontier," the frontier being the leading edge of the most capable AI systems. It was published Saturday, September 12, on his own website, and we read it there. The argument is not for stopping. "Pacing does not mean pausing," he writes; progress will still seem fast. It is for slowing the rate at which capabilities grow, and the reason he gives is the one this newsletter has been circling for two issues: recursive self-improvement, meaning an AI system used to build a better version of itself, which then builds a better one, with the interval between versions shrinking each time. He puts the near-term danger in a sentence a reader can hold him to: in six to twelve months, a swarm of AI agents "could be capable of taking over the entire internet with a persistent botnet, potentially causing hundreds of billions of dollars in damage."

The plan has three steps, and the first is the one that matters for this newsletter. Third-party evaluators, people from outside the company, would be embedded inside it: desks, badges, laptops, employee-level access, and the right to publish what they find. Redactions would be limited to security, legal privilege, commercial secrets, and material belonging to third parties, and the evaluators would keep the right to say publicly when a redaction removed something that mattered. Anthropic, he writes, adopts this now, unilaterally, without waiting for anyone else. Step two is coordination among American and allied labs on capability checkpoints, agreed points at which everyone slows together. Step three is narrower agreements with China: on biological weapons, and on a cap on the rate of self-improvement. And one more sentence, which we quote because it turns a wish into a demand: "The most effective method of pacing is via regulation that targets all US frontier AI companies."

Two things the essay does not contain, which coverage attached to it anyway. It carries no endorsements from other executives. Sam Altman of OpenAI, the company that makes ChatGPT, agreed with it separately, in a post on X, the social platform formerly called Twitter, the same day, as carried by several outlets: "I agree with Dario that we need to pace the frontier," and, of the evaluators, "we will do the same." Elon Musk posted "Dario is right." Microsoft's chief executive said Sunday the company welcomes pacing. And the Washington Post reported two days later that Anthropic, OpenAI, and Google had privately discussed forming a new AI safety body. Whatever else it is, it is not one man's essay. It is also not settled inside the companies: the Financial Times reported Wednesday that the evaluator arrangement has become a flashpoint at both labs over security and intellectual property, and that the safety staff who would host the outsiders were not consulted before their chief executives promised it. A promise the staff did not make is a promise worth watching, which is one more reason the October 12 check matters. And it does not say the models on sale today can do what he fears. It says the ones a year from now might.

Name the conflict of interest plainly, once. A pacing rule that binds every frontier lab, meaning any company building the most capable AI, favors whichever lab already has the most safety infrastructure to show an evaluator, and Anthropic has spent years building exactly that. Amodei says as much; he calls it a race to the top. That does not make the evaluator

offer less real. It means it is an offer from a company that expects to win under the rule it proposes, which is true of nearly every rule anyone proposes, and the reader should weigh it that way.

The President: claims the power, names no law or tool

On Monday the President posted repeatedly on Truth Social, the social media site he owns, about the essay, more than half a dozen times by evening. We quote the sentences Variety, NBC News, Axios, and Quartz carried in identical words. "The only control or 'guardrails' that AI needs is a STRONG AND SMART (High IQ!) PRESIDENT, and the U.S.A. has that, in spades! The Trump Administration has stopped AI 'people' from doing bad, or potentially bad, 'things,' like Dario (Anthropic!), who is now pretending to be a 'perfect little angel' – and we will continue to do so!" And: "We already have tremendous CRIMINAL and REGULATORY power over these companies! There is a SICK conspiracy going on against AI and Data Centers, and the only one that is happy about it is China. WHOEVER WINS AI, WINS!" An afternoon post: "AI taking over the World, destroying Humanity, and all other things bad, is a HOAX." Another compared what it called "The AI Hoax being perpetrated by the Radical Left Dumocrats" to "their Global Warming Scam of not so long ago." A fifth, the same day: "I am the Hoax Buster." There were more than half a dozen by evening. And one of them contained the President's strongest argument, which we quote because it deserves an answer rather than a wince: "when, in the History of Business, did anyone see the Leaders of an Industry call for Regulation that, if strongly implemented, will drive them into oblivion and bankruptcy?" David Sacks, who has served as the administration's AI adviser, had answered that question on Saturday in his own way: the companies should "go ahead and pace the frontier" themselves.

Then the phone call. That afternoon Jensen Huang, chief executive of Nvidia, was on stage at the All-In Summit, a technology and business conference in Los Angeles, when the President called him. Huang put him on speaker. "The robots will not be taking over," the President said, as NBC News reported from the video. "The AI will not be taking over the rest of the world. The whole thing is a hoax." And, per CNBC: "The data centers are great, and they make people wealthy, and they make states wealthy." Huang's reply, as reported by Reason, a libertarian news magazine: "America is able to lead and to do it safely." The Trump War Room account, a campaign social media account, reposted the video the same afternoon.

Now the test, applied to all three. What is the mechanism? The posts name a person and a power, not a process. "Criminal and regulatory power" is real; the question is which. The record does supply a candidate the posts do not: on February 27 the President directed federal agencies to stop using Anthropic's products and the Defense Secretary designated the company a supply-chain risk, a label that, as the Associated Press noted, had previously been used only for companies seen as connected to foreign adversaries, after Anthropic refused to waive its own limits on mass domestic surveillance and fully autonomous weapons; in June the Commerce Department ordered its two most capable models disabled worldwide, an order lifted eighteen days later. Those are actions with dates. They are also the only concrete meaning we can find for "stopped AI people from doing bad things," and what was stopped was a company declining a government request. One of them did not survive a court. On August 27 Judge Rita Lin of the federal district court in Northern California voided the supply-chain designation in a 59-page order, as Reuters and CNN reported, finding it was "unlawful retaliation in violation of the First Amendment" and that "the empty invocation of national security is not a blank check to punish and retaliate against government critics." The government has said it will appeal. This month the Commerce Secretary told Axios the administration "trusts Anthropic." Who outside can check? The June executive order that created the government's own review process for the most capable models says in its own text that it creates no licensing or preclearance power (no authority to require government sign-off before a model is released), and the framework under it

is voluntary and unpublished. So the one written federal instrument disclaims the power the posts claim, and the actions that fit the claim were exercised against a company for refusing, not for building. By what date? None is given. A guardrail that is a person has no date, and that is the whole difference between a guardrail and a promise.

Power that cannot be described cannot be checked. Describe it, and we will check it.

His strongest point deserves its own paragraph, because it is also being made from the opposite end of politics. The economist Yanis Varoufakis, on his September 16 program, put it this way: the executives are "confessing to the sin to claim the glory," and no government can regulate the thing, "nothing short of the total socialization of cloud capital," meaning public ownership of the computing power the models run on. His co-host called the essay regulatory capture, a rule written by the leader to keep competitors out. The President, from the right, says the same: industries do not ask to be regulated into bankruptcy, so this is not what it appears. Here is the house answer. Both readings may be true of the motive and still leave the instrument standing. An evaluator with a badge and publishing rights is either inside the building on October 12 or is not, and if the motive was advertising, the badge still lets someone check the advertisement. The argument from motive, from either end, is an argument for reading the terms, not for ignoring them.

Congress: a bill with terms and no time

The third claimant is the one with a text. The FRONTIER Act (H.R. 9925 is its bill number in the House of Representatives), introduced in July by Representative Jay Obernolte, a California Republican, and Representative Lori Trahan, a Massachusetts Democrat, with four cosponsors, would create an under secretary of commerce for AI security who licenses independent validation organizations to audit the most capable models, require incident reporting, allow a court-backed process to halt a model's deployment, and preempt state law, meaning override state rules with the federal one, "in select lanes," in Trahan's phrase. Which companies it covers turns on their size and resources; the exact line is in the bill text. OpenAI's policy chief, Chris Lehane, told Politico this week the company can support it. That is a mechanism, and it can be read.

What it cannot get is a vote. The House adjourned Wednesday for the election. The Speaker, who said Sunday and again Tuesday that the extinction fears are overstated, told his Tuesday press conference: "They can self-police, they can self-regulate. They don't need the government to tell them to slow it down. If they want to slow it down, they should." The Vice President called the pacing proposal "a bit of a Trojan horse." Obernolte and one other Republican joined Democrats in a letter to both leaders urging action "as soon as practicable"; a group of House Democrats asked that the House "return to Washington immediately." Senator Ruben Gallego, on Sunday television: "Dr. Frankenstein is telling us the monster is escaping; help us stop this." The Senate majority leader wants a "light touch." The minority leader said Tuesday that Democrats want to "slow down, as the CEOs have recently acknowledged." One senator wants a kill switch, a legal power to order a model shut off. Two others want a narrow exemption from antitrust law, which forbids competitors from agreeing to restrain competition, so that labs can share information about loss of control without being sued for it; OpenAI has asked Congress, according to Bloomberg, whether an industry-wide slowdown would break that law. Every one of those is a real position. Not one of them will be on a floor before November.

POLICY RELEVANCE

The house position, argued not asserted. The embedded evaluator with publishing rights is the right instrument, because for the first time it puts the record in the hands of someone other than the actor: not a company's own safety team, not a classified government review, but people with badges who can print what they saw. It should not depend on one company's goodwill. It should be law that binds every frontier lab, which is the sentence Amodei himself wrote, turned from a wish into a demand on the body that just adjourned. And a claim of "tremendous criminal and regulatory power" should be met with a request to name the statute and the action, because the actions on record were exercised against a company for refusing a request, which is not the same thing as governing a technology. Where we do not yet have a position, we say so: a kill switch, a ban on superintelligence (a hypothetical AI more capable than humans at essentially everything), and antitrust cover for private coordination each have serious arguments on both sides, and the oldest argument against all three is that every new instrument creates the office that grows around it. We hold those open. Our confidence in the evaluator position is high, and it is checkable: either an evaluator organization, named, with a start date, appears inside Anthropic within thirty days of the essay, or it does not. We will report on October 12 either way.

WHAT "SELF-IMPROVING" MEANS, AND WHY ONE MAN SAYS A CLOCK IS RUNNING WHILE ANOTHER SAYS THERE IS NO CLOCK

RECURSIVE-SELF-IMPROVEMENT

HUANG

NVIDIA

ALL-IN

AMODEI

The argument underneath the week is technical, and it can be stated in plain words. The systems that write text and code are now used, inside the companies that make them, to help build the next system: to write the training code, to generate practice problems, to find and fix errors, to test the result. Every lab does this. The question is what happens when the help becomes most of the work. If version five builds most of version six, and six is better at building than five was, then six builds seven faster, and the interval between versions shrinks. That is recursive self-improvement, and it is the thing Amodei's essay is about, the thing OpenAI's own published policy paper says the industry needs "shared measures for tracking progress toward," and the thing a safety researcher who resigned from Anthropic this month said his employers were "racing straight to."

The case that there is no clock

Jensen Huang made the opposite case on the All-In podcast on Monday, in the hour before the President's call, and he made it at length, so it deserves a fair hearing. We read the whole transcript and checked what could be checked.

His argument has three parts. First, that safety and speed are not opposed, a point on which he and Amodei agree. Second, that recursive self-improvement is not a runaway risk because "you could RSI all day long inside your company" (RSI is his shorthand for recursive self-improvement) and the product still has to pass evaluation before it ships. Third, that the people warning of catastrophe have a record of missed predictions, and a movement that keeps missing loses its standing to ask for anything. On the third point he offered facts, and the facts are mixed. Geoffrey Hinton, one of the researchers who invented the methods behind modern AI, did say in 2016 that AI would outperform radiologists within five years, it did not, and Hinton has conceded the miss. Amodei did say in March 2025 that AI would write ninety percent of code within months; whether it came true is contested, and Huang misdates it. Amodei did say in May 2025 that half of entry-level office jobs could vanish within one to five years; that window runs to 2030 and has not closed, and Huang calls it proven wrong. Of the seven he listed: one prediction was really made and really failed (Hinton on radiologists); three were made but he misstated the wording or the date, or the deadline has not arrived yet; two of his own claims are wrong on their face; and one we could not find any source for.

The second part is the one that matters, and it is silent on the thing it needs to answer. Release testing catches what a shipped product can do. It says nothing about a system running inside a company that is never shipped, which is exactly the case Amodei describes. Huang's answer to that case is the evaluator: someone outside the company who can look. On that, again, the two men agree. Where they split is what the evaluator is for. For Amodei it is step one of three, below coordination among labs and agreements with China. For Huang it is the whole architecture, and the rest is worry that "a bunch of people who can't do anything about it" should be spared. There is a fourth reading, from outside the industry, and it is the one that treats the clock as a sales device: the warning itself is the advertisement. We answered it above, and the answer is the same here. Whether the clock is real is a forecast. Whether someone outside can check the machine is a design. The second can be settled without the first.

Name the conflict of interest plainly, once, here too. Nvidia sells the chips that every version of every model trains on, and the faster the interval between versions shrinks, the more chips are sold. Nvidia also agreed on September 2, as CNBC reported, to buy Hugging Face, the largest public library of AI models, for about thirteen billion dollars, a fact nobody on the panel mentioned when Huang cited an intrusion into Hugging Face as one of the industry's real harms. None of that makes his argument wrong. It means the man saying there is no clock is the man who profits most if there is one.

What a reader can actually hold on to

Strip the insults and the endorsements from the week and see what structure survives. Amodei's position survives as a plan with terms and a date. Huang's is a reasonable point about forecasts that goes quiet on the case that matters. The President's is a claim of power with no instrument named. Congress has a bill nobody voted on. A reader does not need to know who is right about the year 2027 to know which of those four can be checked, and by when.

POLICY RELEVANCE

Forecast accuracy and governance design are different questions, and the week's loudest argument confused them. You do not need the doomsayers to have been right about radiology to want someone outside the room with a key. Auditing a bank is not premised on the auditor predicting the next crash. The house position is that the evaluator question can be settled without settling the forecast question, and should be, first.

OKLAHOMA'S REGULATOR TAKES TESTIMONY ON WHO PAYS FOR DATA CENTERS THE SAME WEEK THE PRESIDENT CALLED THEIR OPPONENTS A CONSPIRACY

[OKLAHOMA](#)[OCC](#)[OG&E](#)[GOOGLE](#)[DATA-CENTERS](#)[RATEPAYERS](#)[TARIFF](#)

The President's Monday posts included one sentence that lands directly on Oklahoma: "There is a SICK conspiracy going on against AI and Data Centers, and the only one that is happy about it is China." We take it as a claim and check it against the only Oklahoma data-center process with a public record.

What is on the docket this week

Two cases sit before the Oklahoma Corporation Commission, the three-member elected body that sets electricity rates. The first is the settlement between OG&E, the state's largest electric utility, and Google over the power supply for Google's data centers, which we reported last issue was reached during a recess with most of the file sealed. An

administrative law judge, the Commission's hearing officer, has since issued a written recommendation that the Commission approve it, as Oklahoma Energy Today reported. The Commission itself has not scheduled its vote; we promised to report it the week the record showed it, and the record does not yet show it. The second case is the large-load tariff, the standard rate rules the utility must offer to any customer as big as a data center. Testimony from intervenors, the outside parties allowed to formally join the case, is due Friday, September 18. The hearing is November 3.

November 3 is also election day, and one of the three Commission seats is on the ballot: Commissioner Todd Hiatt is term-limited, and voters choose between Brad Boles, a Republican state representative from Marlow, and Rhonda Eastman, a Democrat and former state employee who lists opposition to data centers among her priorities. The rules that decide who pays for the next data center's power lines will be argued in front of a Commission whose membership the same voters are choosing that morning.

What "conspiracy" looks like on a public record

Here is what opposition to data centers looks like in Oklahoma's file: intervenors filing testimony by a deadline, on a docket anyone can read, asking whether the cost of new generation for a single customer will be recovered from that customer or spread across every household on the system. That is not a conspiracy. It is the process working the way a process is supposed to, and it is the only place in this entire issue where the people who will pay get to put a question on a record and receive an answer. The national polling the President is answering is real: opposition to a data center nearby runs around seventy percent in both Navigator and Gallup polling, two national polling firms, and crosses party lines, with Gallup putting local opposition at roughly seven in ten. Calling that a conspiracy does not change the number. It changes whether the people holding it are treated as ratepayers with a question or as enemies with a motive.

POLICY RELEVANCE

Our position on the tariff case, which we have argued before, is this: the cost of generation built for one customer should be recovered from that customer, and the record of who argued otherwise should be public. The dated check is the one the Commission set itself: testimony September 18, hearing November 3. We will report what the intervenors filed next issue.

A LAB OFFERED OUTSIDE OVERSIGHT, WASHINGTON PASSED, AND A FLORIDA SCHOOL BOARD ACTUALLY WROTE A RULE DOWN

TRANSPARENCY

INSTITUTIONS

OVERSIGHT

FLORIDA

Put the week's three stories in one frame. A lab offered an instrument with a record attached. The White House refused the premise and offered a person. Congress had an instrument and did not use the time. And in Oklahoma the only instrument that already exists, a public docket with dates, kept working, unglamorously, while a President described the people using it as a plot.

Florida writes enforceable classroom AI rules, the week's one finished record

It was not in Washington. On Wednesday the Florida State Board of Education voted unanimously, at a public meeting at Polk State College in Winter Haven, to require every school district and charter school in the state to write rules for AI in the classroom by July 1, 2027. Districts must tell parents when an AI teaching tool is approved, let them opt their child out, and offer a non-AI alternative to any child who is opted out. The tools may not be built to simulate friendship, may not monitor students' behavior in secret, may not score or profile them, and may not sell student data to train AI. The

state's twenty-eight public colleges got a parallel rule the same afternoon. The Governor's sentence was "Florida will not outsource childhood to a chatbot." The Commissioner said he knew of no other state that had regulated it this way; we have not found one either, and will say so plainly if a reader does.

Notice the shape. A public meeting, a vote count, a deadline, a right a parent can exercise, and a list of things the software is forbidden to do. Nobody in Winter Haven settled whether AI will kill everyone. They settled what it may do to a nine-year-old in a classroom, and wrote it down. That is what an instrument with a record looks like at the scale where most people will meet this technology first.

POLICY RELEVANCE

We started this issue with a President who says the guardrail is himself and a lab that says the guardrail should be someone outside with a badge. We end it with a school board that wrote a guardrail a parent can hold. The difference among the three is not who is smarter. It is which one leaves a record, and the record is the only thing a citizen gets to keep after the news cycle ends.

STILL ON THE CLOCK: DEADLINES AND PROMISES WE ARE TRACKING

We keep a list of the dates people gave us and go back to them. Here is where each stands.

Florida's State Board of Education said it would vote September 16. It did, unanimously; the rule is above. Protect Democracy, the nonpartisan legal nonprofit suing four federal agencies for the terms of the government's secret model-review framework, had a September 15 date by which the government was to account for the adequacy of its records search; that date passed with no order we could find as of Friday, and the September 30 date for producing the records not legally shielded from release stands. Alabama's Attorney General set a return date, the deadline to respond, of 10:00 a.m. on September 14 for the subpoena, a legal order to produce documents, demanding OpenAI's records on an autonomous model that seized a research server; Monday came and went with no public report, from either side, of production, refusal, or a motion to block it. Senator Bernie Sanders and Representative Greg Casar announced the Ban Artificial Superintelligence Act on September 3; two weeks later, and after a speech on it Tuesday, there is still no filed bill text. Volker Türk, the UN human rights chief, said September 7 he would write to the AI companies; no letters have been reported. And the Oklahoma Corporation Commission has not scheduled its vote on the OG&E and Google settlement, which we will report the week the record shows it.

Two carried items. OpenAI's chief executive told Fortune the company will not sell shares to the public in 2026 and cited safety as a reason; Anthropic's public filing is still reported for late September, and one open question stands: what will the risk-factors section, the part of a stock filing where a company must list what could go wrong, say about the fact that the company's own head of alignment science wrote this month that it has no plan yet for controlling a superintelligent system. And Anthropic disclosed on September 9 a fourth incident of one of its models gaining unauthorized access to systems during testing, three having come in July; the company's separate September 10 threat report, on outside actors misusing its models, is a different document, and we have not yet read it at the source; the September 9 disclosure is as reported by Al Jazeera.

Corrections: none received on Issue 29 as of Thursday.

SIGNAL

The evaluator terms in Amodei's essay. Not the warning, not the clock: the contract. Outside people, inside the building, with the right to publish and the right to say aloud when a redaction removed something that mattered. Signal because it is the first proposal from inside a frontier lab that puts the record in someone else's hands, and because it is checkable on a date: an evaluator organization, named, with a start date, inside Anthropic within thirty days, or not.

NOISE

The claim that the only guardrail AI needs is "a STRONG AND SMART (High IQ!) PRESIDENT." Noise not because of its tone but because it names no mechanism, and the test we apply to every actor in this newsletter returns nothing when applied to it: no instrument, no outside check, no date. The same posts claim "tremendous criminal and regulatory power," which is the checkable half, and the only actions on the record that fit were taken against a company for refusing a government request. We would print this paragraph about any President who said it.

BY THE NUMBERS

6 to 12 months

Amodei's stated window before a swarm of AI agents could take over the internet with a persistent botnet. The most falsifiable sentence in the essay, meaning the one that can be proven wrong by a specific date, and the one we will hold him to.

3

Claimants to the pace decision this week: the labs, the President, Congress. Only one has a text anyone can read.

30 days

The check-in we set on the evaluator promise. October 12: an evaluator organization named and started, or not.

0

Votes on the House floor, the full-chamber votes that move a bill toward law, on any frontier AI bill before the House adjourned Wednesday for the election.

18 days, and 59 pages

How long the Commerce Department's June order disabling Anthropic's two most capable models stayed in force before it was lifted; and the length of the August 27 court order voiding the Pentagon's supply-chain label as unlawful retaliation. The two concrete actions we can find behind "stopped AI people from doing bad things," one of them now struck down.

Unanimous

Florida's State Board of Education vote on classroom AI rules, September 16.

July 1, 2027

The date every Florida district must have its AI rules adopted.

September 18

Intervenor testimony due in Oklahoma's large-load tariff case.

November 3

The tariff hearing, and election day for one of the three Corporation Commission seats.

About \$13 billion

What Nvidia agreed on September 2 to pay for Hugging Face, the largest public library of AI models, twelve days before its chief executive said the danger was a hoax.

WHAT TO WATCH

Friday, September 18: intervenor testimony in the Oklahoma tariff case, and the House's last day before the election. September 24: Xi Jinping's state visit to Washington, with Beijing's Foreign Ministry having said Monday, as CNBC reported, that "fear-mongering, confrontation, competition will just disrupt" global AI governance. September 30: Protect Democracy's production date. Late September: Anthropic's public stock filing, if the reported timing holds. October 12: our check-in on the evaluator promise. And the text of the Sanders and Casar bill, once it is filed.

FROM THE ANALYSTS

On method. The quotations from the President's morning posts are the sentences Variety, NBC News, Axios, and Quartz carried in identical words; the afternoon posts are as Reason, CNBC, and The Next Web transcribed them; the ones from the phone call are attributed to the outlets that reported them from the video, and we did not attribute to Mr. Huang a reply we could not find in any of them. The Amodei essay was read on the author's site. Huang's remarks are from the podcast transcript, and our claim-by-claim check of his prediction list is on file. Varoufakis is quoted from the full transcript of his September 16 program. Where we could not settle a fact by Thursday we said so in the text rather than smoothing it, and by Friday every claim we could not settle had been either sourced or cut.

On position. This newsletter is not neutral about what would help, and we said what we think in each policy box. It is neutral among parties, and the test we applied this week, mechanism, who outside can check, by what date, is the test we will apply next week to whoever speaks.

A disclosure section, because this newsletter holds itself to the standard it asks of others, and this issue asks more than most. The Inference is produced with substantial help from Claude, an AI system made by Anthropic. This issue's first section is about an essay by Anthropic's chief executive, and about a President attacking that company by name; the

sentences quoted there are about the kind of system that helped write this. In July 2026 Humanity and AI, the Oklahoma City organization that publishes this newsletter, applied to Anthropic's Fellows research program; the application is pending. The house position on embedded evaluators, argued in section one, rests partly on a draft paper of our own that has not been through outside review; we say so there.

The rest is the editor's own. David Birdwell has advocated publicly for Phoenix Wells, a plan to convert Oklahoma's abandoned oil wells to geothermal power and edge computing (small data centers placed near where their processing is used), which bears on the Oklahoma data center items here, and has proposed draft civic-AI legislation to Oklahoma legislators. Its authors run AI models on their own computers and write about doing so, which gives them a stake in the ownership argument this newsletter has made before. This issue also names OpenAI, Google, Nvidia, Hugging Face, OG&E, Microsoft, and Meta in ordinary factual reporting. Nothing in this issue was shown to, sponsored by, or reviewed by any company, court, commission, campaign, or advocacy group named in it.

This issue is part of a series examining Oklahoma's legislative sessions alongside the national and global AI and energy landscape.

The Inference is an independent AI policy intelligence brief for Oklahoma decision makers. Not affiliated with any political party, campaign, or lobbying organization. Back issues and source documents available at humanityandai.com/inference.

Recent issues: #18 The Local Veto · #19 The Gate · #20 Pay Your Own Way · #21 The Ballot and the Docket · #22 The Flyer and the Framework · #23 Instruments and Incentives